

1. My Machine:

I ran this program on flip in linux from a windows 10 laptop. My main function is contained in a file called project5.cpp

I wrote a script to automate the program in a file called runProj5.py . It can be compiled and executed by typing:

```
python3 runProj5.py
```

2. My performance results:

Array multiplication:

ARRAYSIZE	C++ Timing (seconds)		SIMD Timing (seconds)		Performance (MegaMults/Sec)		Speedup
	C++ Peak	C++ Average	SIMD Peak	SIMD Average	C++ Peak	SIMD Peak	
1000	0.00000827	0.00000839	0.00000103	0.0000011	120.9189843	970.8738	8.029126214
10000	0.00008271	0.00008352	0.00000992	0.00001001	120.9043646	1008.065	8.337701613
30000	0.00024833	0.00025103	0.0000359	0.00003631	120.8069907	835.6546	6.917270195
50000	0.00041362	0.0004166	0.00006113	0.00006177	120.8839031	817.929	6.766235891
70000	0.0005782	0.00058277	0.00008646	0.00008724	121.0653753	809.6229	6.687485542
100000	0.00082887	0.00083428	0.00012251	0.00012349	120.6461809	816.2599	6.76573341
200000	0.00153594	0.00154082	0.00022853	0.00023855	130.2134198	875.1586	6.720955673
300000	0.00230479	0.0023821	0.00038184	0.00038431	130.1637025	785.6694	6.036009847
400000	0.00173617	0.00319755	0.00027334	0.00027502	230.3921851	1463.379	6.351686544
500000	0.00217172	0.00217961	0.0003424	0.00034451	230.2322583	1460.28	6.342640187
600000	0.00260117	0.0026217	0.00041629	0.00078643	230.6654313	1441.303	6.248456605
700000	0.0030495	0.00514608	0.0004824	0.00049053	229.5458272	1451.078	6.321517413
800000	0.00347935	0.00468217	0.00055887	0.00056433	229.9280038	1431.46	6.225687548
900000	0.00392048	0.00512529	0.00063933	0.00064882	229.5637269	1407.724	6.132169615
1000000	0.00437696	0.00581575	0.00075662	0.00079198	228.4690744	1321.667	5.784885411
2000000	0.0089358	0.01014341	0.00228457	0.00231108	223.8187963	875.4383	3.91137063
3000000	0.01343896	0.01472484	0.00339018	0.00341478	223.2315596	884.9088	3.964084503
4000000	0.01792593	0.01916757	0.00447789	0.00451504	223.1404452	893.2779	4.003209101
5000000	0.02235789	0.02651216	0.00548688	0.00665285	223.634699	911.2647	4.074791138
6000000	0.02684699	0.02942486	0.00655254	0.00658659	223.4887412	915.6754	4.097188266
7000000	0.03125881	0.03261286	0.00853321	0.00869485	223.9368677	820.3244	3.663194741


8000000	0.03589608	0.03609282	0.0088829	0.00953514	222.8655608	900.6068	4.041031645
9000000	0.04030456	0.04144555	0.01029528	0.0105207	223.2997954	874.187	3.914858071
10000000	0.0448177	0.04497463	0.01110521	0.0111834	223.126131	900.4782	4.03573638
15000000	0.06721246	0.0673972	0.01730253	0.01752695	223.1729057	866.9252	3.884545208
20000000	0.08958104	0.08985086	0.02246897	0.0225907	223.2615294	890.1165	3.986877903
25000000	0.11202863	0.11261517	0.02857529	0.02875641	223.1572411	874.8818	3.920472198
32000000	0.14209378	0.14286236	0.03676769	0.03817045	225.2033833	870.3294	3.864637131

Array multiplication-reduction:


ARRAYSIZE	C++ Timing (seconds)		SIMD Timing (seconds)		Performance (MegaMults/Sec)		Speedup
	C++ Peak	C++ Average	SIMD Peak	SIMD Average	C++ Peak	SIMD Peak	
1000	0.00000429	0.00000432	0.00000054	0.00000055	233.1002331	1851.851852	7.944444444
10000	0.00004237	0.0000427	0.00000521	0.00000523	236.0160491	1919.385797	8.13243762
30000	0.00012709	0.00012803	0.00001577	0.00001584	236.0531907	1902.346227	8.058972733
50000	0.00021216	0.00021436	0.00002661	0.00002679	235.6711916	1878.99286	7.972942503
70000	0.00029697	0.00029836	0.00003732	0.00003751	235.7140452	1875.669882	7.957395498
100000	0.00042421	0.00042628	0.00005334	0.00005358	235.7323024	1874.765654	7.952943382
200000	0.00084847	0.00087856	0.00010667	0.00010909	235.7184108	1874.941408	7.954157683
300000	0.00127578	0.00128505	0.0001596	0.0001602	235.150261	1879.699248	7.993609023
400000	0.00169797	0.00171308	0.00021346	0.00021451	235.5754224	1873.887379	7.954511384
500000	0.00212247	0.00212963	0.00026682	0.0002679	235.57459	1873.922495	7.954688554
600000	0.0025467	0.00255698	0.00031772	0.00031981	235.5990105	1888.455244	8.015548282
700000	0.00297323	0.00298359	0.00037238	0.00037498	235.4341911	1879.800204	7.984397658
800000	0.00339908	0.00341315	0.00042518	0.00042787	235.3578027	1881.556047	7.99444941
900000	0.00382337	0.00394036	0.00049773	0.00050041	235.3944295	1808.20927	7.68161453
1000000	0.00443653	0.00444426	0.00055705	0.00055896	225.4013835	1795.17099	7.964329952
2000000	0.00868328	0.00884171	0.00133465	0.0013507	230.3277103	1498.520211	6.50603529
3000000	0.01310586	0.01468074	0.00213912	0.00216112	228.9052378	1402.445866	6.126753057
4000000	0.01745863	0.01845483	0.00283087	0.00286002	229.1130518	1412.993179	6.167231275
5000000	0.0218606	0.02194306	0.00353476	0.00364812	228.721993	1414.523193	6.18446514
6000000	0.02620366	0.02760191	0.00432862	0.00436735	228.9756469	1386.12306	6.053582897
7000000	0.03056759	0.03066942	0.00507439	0.00514658	229.0007161	1379.476154	6.023894498
8000000	0.03493791	0.03610444	0.00572426	0.00620947	228.9776349	1397.560558	6.103480625
9000000	0.03927932	0.04156572	0.00650238	0.00672532	229.1282028	1384.108588	6.04076046
10000000	0.04365861	0.04464442	0.00711545	0.00716805	229.0498942	1405.392491	6.135748266
15000000	0.06554296	0.07048181	0.01093265	0.01145187	228.857531	1372.036972	5.995157624
20000000	0.08731087	0.08757948	0.01421706	0.01435095	229.0665527	1406.76061	6.141274638
25000000	0.10917367	0.1096608	0.01815522	0.01875734	228.9929431	1377.014434	6.013348778
32000000	0.1380887	0.13835394	0.02259804	0.02324053	231.7351094	1416.052012	6.110649419

Graphs of Performance:

Array Size vs Performance (Multiplication)

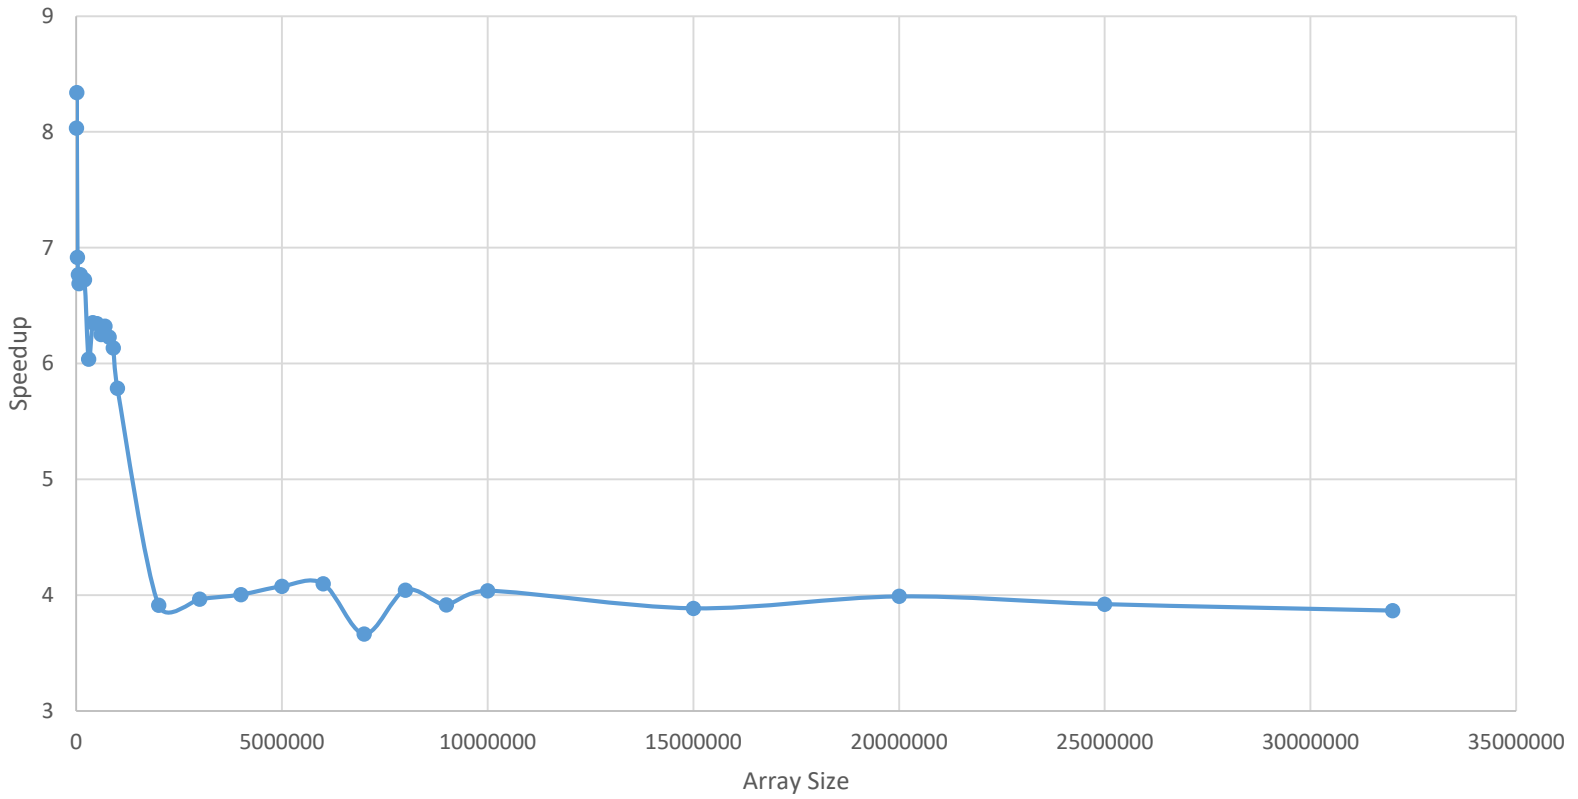


Array Size vs Performance (Multiplication-Reduction)




Graphs of Speedup:

Array Size vs Speedup (Multiplication)



Array Size vs Speedup (Multiplication-Reduction)



3. Speedup Patterns:

For array multiplication, the speedup (of SIMD to non-SIMD) starts around 8 for the smallest array sizes, then there is a huge drop to around 4 and then it remains consistently around 4 as array size increases.

For array multiplication-reduction, the speedup (of SIMD to non-SIMD) also starts around 8 for the smallest array sizes and then there is a drop to around 6 and it remains consistently around 6 as array size increases.

4. Consistency

When the array size is between 1,000-200,000 for array multiplication (and 1,000-300,000 for array multiplication-reduction), we see a dramatic decrease in the speedup values (from 8 to 4, and 8 to 6 respectively). Within this range, as the array size increases, speedup drastically decreases.

Once the array size reaches 200000 for array multiplication (and 300000 for array multiplication-reduction), the speedup values are very consistent. For multiplication, the speedup stays around 4 and for multiplication-reduction, the speedup stays around 6.

Regardless of where we look at these graphs, we can see that the SIMD implementation is still a considerable improvement from the C++ implementation for any array size.

5. Explanation of Consistency

I believe that the initial dramatic drops in speedup are due to cache misses. SIMD requires 4 times the space in cache as non-SIMD and the CPU is trying to fetch data from the cache that isn't being filled in time. Initially, the cache is filled with values we need which is why we have the most dramatic speedup initially. Then, as array size grows, we need to fill the cache line with more values but they are needed faster than they can be delivered.

It makes sense that the larger array sizes are consistent because the SIMD code is able to perform 4x the number of mathematical operations at a time as the non-SIMD code.

6. Array multiplication speedup != 4

Speedup for array multiplication could be less than 4 because of cache misses due to temporal coherence and the program requiring 3 arrays to store each of the 2 numbers to be multiplied as well as the product in a third array.

Speedup could be more than 4 if the values needed are already on the cache line and can be accessed very quickly.

7. Array multiplication-reduction speedup != 4

Speedup for array multiplication-reduction could be less than 4 because of cache misses due to temporal coherence and the program requiring 2 arrays to store each of the 2 numbers to be multiplied then added.

Speedup could be more than 4 because when performing the multiplication-reduction in C++, the sum is stored in a local variable, which is stored to memory via cache. However, when performing the same operation with SIMD, the sum is stored in another xmm register which is much faster to access.